

gawk(1) — Linux manual page

[NAME](#) | [SYNOPSIS](#) | [README FIRST](#) | [DESCRIPTION](#) | [PREFACE](#) | [OPTION FORMAT](#) | [OPTIONS](#) | [AWK PROGRAM EXECUTION](#) | [VARIABLES, RECORDS AND FIELDS](#) | [PATTERNS AND ACTIONS](#) | [USER-DEFINED FUNCTIONS](#) | [DYNAMICALLY LOADING NEW FUNCTIONS](#) | [SIGNALS](#) | [INTERNATIONALIZATION](#) | [GNU EXTENSIONS](#) | [ENVIRONMENT VARIABLES](#) | [EXIT STATUS](#) | [VERSION INFORMATION](#) | [AUTHORS](#) | [BUG REPORTS AND QUESTIONS](#) | [BUGS](#) | [SEE ALSO](#) | [EXAMPLES](#) | [COPYING PERMISSIONS](#) | [COLOPHON](#)

GAWK(1)

Utility Commands

GAWK(1)

NAME [top](#)

gawk – pattern scanning and processing language

SYNOPSIS [top](#)

```
gawk [ POSIX or GNU style options ] -f program-file [ -- ] file
...
gawk [ POSIX or GNU style options ] [ -- ] program-text file ...
```

README FIRST [top](#)

This manual page is provided as a *courtesy*. Please note that the One Source Of Truth for *gawk* is the Texinfo manual, available online in several formats at <https://www.gnu.org/software/gawk/manual> . It may also be installed in the Info subsystem on your system, and available therefore via the *info*(1) command.

In the case of any contradiction between the Texinfo manual and this manual page, the Texinfo manual should be considered to be authoritative.

Note that the Info documentation, HTML web page(s), and PDF version of the manual all have the same content, since they are all generated from the same Texinfo source file.

Furthermore, the *gawk* maintainer reserves the right in the future to remove this manual page from the distribution or to severely truncate the content in order to ease his maintenance burden.

DESCRIPTION [top](#)

Gawk is the GNU Project's implementation of the AWK programming language. It conforms to the definition of the language in the POSIX 1003.1 standard. This version in turn is based on the description in *The AWK Programming Language*, by Aho, Kernighan, and Weinberger. *Gawk* provides the additional features found in the current version of Brian Kernighan's *awk* and numerous GNU-specific extensions.

The command line consists of options to *gawk* itself, the AWK program text (if not supplied via the *-f* or *--include* options),

and values to be made available in the `ARGC` and `ARGV` pre-defined AWK variables.

PREFACE

[top](#)

This manual page is intentionally as terse as possible. Full details are provided in *GAWK: Effective AWK Programming*, and you should look there for the full story on any specific feature. Where possible, links to the online version of the manual are provided.

OPTION FORMAT

[top](#)

Gawk options may be either traditional POSIX-style one letter options, or GNU-style long options. POSIX options start with a single “-”, while long options start with “--”. Long options are provided for both GNU-specific features and for POSIX-mandated features.

Gawk-specific options are typically used in long-option form. Arguments to long options are either joined with the option by an = sign, with no intervening spaces, or they may be provided in the next command line argument. Long options may be abbreviated, as long as the abbreviation remains unique.

Additionally, every long option has a corresponding short option, so that the option's functionality may be used from within `#!` executable scripts.

OPTIONS

[top](#)

Gawk accepts the following options. Standard options are listed first, followed by options for *gawk* extensions, listed alphabetically by short option.

- `-f program-file, --file program-file`
Read the AWK program source from the file *program-file*, instead of from the first command line argument. Multiple `-f` options may be used. Files read with `-f` are treated as if they begin with an implicit `@namespace "awk"` statement.
- `-F fs, --field-separator fs`
Use *fs* for the input field separator (the value of the `FS` predefined variable).
- `-v var=val, --assign var=val`
Assign the value *val* to the variable *var*, before execution of the program begins. Such variable values are available to the `BEGIN` rule of an AWK program.
- `-b, --characters-as-bytes`
Treat all input data as single-byte characters. The `--posix` option overrides this one.
- `-c, --traditional`
Run in *compatibility* mode. In compatibility mode, *gawk* behaves identically to Brian Kernighan's *awk*; none of the GNU-specific extensions are recognized.
- `-C, --copyright`
Print the short version of the GNU copyright information message on the standard output and exit successfully.
- `-d[file], --dump-variables[=file]`
Print a sorted list of global variables, their types and

final values to *file*. The default file is **awkvars.out** in the current directory.

-D[*file*], --debug[=*file*]

Enable debugging of AWK programs. By default, the debugger reads commands interactively from the keyboard (standard input). The optional *file* argument specifies a file with a list of commands for the debugger to execute non-interactively.

In this mode of execution, *gawk* loads the AWK source code and then prompts for debugging commands. *Gawk* can only debug AWK program source provided with the **-f** and **--include** options. The debugger is documented in *GAWK: Effective AWK Programming*; see

https://www.gnu.org/software/gawk/manual/html_node/Debugger.html#Debugger .

-e *program-text*, --source *program-text*

Use *program-text* as AWK program source code. Each argument supplied via **-e** is treated as if it begins with an implicit **@namespace "awk"** statement.

-E *file*, --exec *file*

Similar to **-f**, however, this option is the last one processed. This should be used with **#!** scripts, particularly for CGI applications, to avoid passing in options or source code (!) on the command line from a URL. This option disables command-line variable assignments.

-g, --gen-pot

Scan and parse the AWK program, and generate a GNU **.pot** (Portable Object Template) format file on standard output with entries for all localizable strings in the program. The program itself is not executed.

-h, --help

Print a relatively short summary of the available options on the standard output. Per the *GNU Coding Standards*, these options cause an immediate, successful exit.

-i *include-file*, --include *include-file*

Load an awk source library. This searches for the library using the **AWKPATH** environment variable. If the initial search fails, another attempt will be made after appending the **.awk** suffix. The file will be loaded only once (i.e., duplicates are eliminated), and the code does not constitute the main program source. Files read with **--include** are treated as if they begin with an implicit **@namespace "awk"** statement.

-l, --trace

Print the internal byte code names as they are executed when running the program. The trace is printed to standard error. Each "op code" is preceded by a **+** sign in the output.

-k, --csv

Enable CSV special processing. See **Comma Separated Values**, below, for more detail.

-l *lib*, --load *lib*

Load a *gawk* extension from the shared library *lib*. This searches for the library using the **AWKLIBPATH** environment variable. If the initial search fails, another attempt will be made after appending the default shared library suffix for the platform. The library initialization routine is expected to be named **dl_load()**.

-L [*value*], --lint[=*value*]

Provide warnings about constructs that are dubious or non-

portable to other AWK implementations. See https://www.gnu.org/software/gawk/manual/html_node/Options.html#Options for the list of possible values for *value*.

-M, --bignum

Force arbitrary precision arithmetic on numbers. This option has no effect if *gawk* is not compiled to use the GNU MPFR and GMP libraries. (In such a case, *gawk* issues a warning.)

NOTE: This feature is *on parole*. The primary *gawk* maintainer is no longer supporting it, although there is a member of the development team who is. If this situation changes, the feature will be removed from *gawk*.

-n, --non-decimal-data

Recognize octal and hexadecimal values in input data. *Use this option with great caution!*

-N, --use-lc-numeric

Force *gawk* to use the locale's decimal point character when parsing input data.

-o[*file*], --pretty-print[=*file*]

Output a pretty printed version of the program to *file*. The default file is **awkprof.out** in the current directory. This option implies **--no-optimize**.

-O, --optimize

Enable *gawk*'s default optimizations upon the internal representation of the program. This option is on by default.

-p[*prof-file*], --profile[=*prof-file*]

Start a profiling session, and send the profiling data to *prof-file*. The default is **awkprof.out** in the current directory. The profile contains execution counts of each statement in the program in the left margin and function call counts for each user-defined function. *Gawk* runs more slowly in this mode. This option implies **--no-optimize**.

-P, --posix

This turns on *compatibility* mode, and disables a number of common extensions.

-r, --re-interval

Enable the use of *interval expressions* in regular expression matching. Interval expressions are enabled by default, but this option remains for backwards compatibility.

-s, --no-optimize

Disable *gawk*'s default optimizations upon the internal representation of the program.

-S, --sandbox

Run *gawk* in sandbox mode, disabling the **system()** function, input redirection with **getline**, output redirection with **print** and **printf**, and loading dynamic extensions. Command execution (through pipelines) is also disabled.

-t, --lint-old

Provide warnings about constructs that are not portable to the original version of UNIX *awk*.

-V, --version

Print version information for this particular copy of *gawk* on the standard output. This is useful when reporting bugs. Per the *GNU Coding Standards*, these options cause an

immediate, successful exit.

- Signal the end of options. This is useful to allow further arguments to the AWK program itself to start with a “-”.

In compatibility mode, any other options are flagged as invalid, but are otherwise ignored. In normal operation, as long as program text has been supplied, unknown options are passed on to the AWK program in the `ARGV` array for processing.

For POSIX compatibility, the `-W` option may be used, followed by the name of a long option.

AWK PROGRAM EXECUTION [top](#)

An AWK program consists of a sequence of optional directives, pattern-action statements, and optional function definitions.

```
@include "filename"
@nsinclude "filename"
@load "filename"
@namespace "name"
pattern { action statements }
function name(parameter list) { statements }
```

Gawk first reads the program source from the *program-file(s)* if specified, from arguments to `--source`, or from the first non-option argument on the command line. The `-f` and `--source` options may be used multiple times on the command line. *Gawk* reads the program text as if all the *program-files* and command line source texts had been concatenated together.

In addition, lines beginning with `@include` may be used to include other source files into your program. This is equivalent to using the `--include` option.

Lines beginning with `@nsinclude` also include files, but with two differences. First, the namespace in the included file is *not* reset to the default of `awk`, and such files may be included multiple times.

Lines beginning with `@load` may be used to load extension functions into your program. This is equivalent to using the `--load` option.

The environment variable `AWKPATH` specifies a search path to use when finding source files named with the `-f` and `--include` options. If this variable does not exist, the default path is `./usr/local/share/awk`. (The actual directory may vary, depending upon how *gawk* was built and installed.) If a file name given to the `-f` option contains a “/” character, no path search is performed.

The environment variable `AWKLIBPATH` specifies a search path to use when finding source files named with the `--load` option. If this variable does not exist, the default path is `/usr/local/lib/gawk`. (The actual directory may vary, depending upon how *gawk* was built and installed.)

Gawk executes AWK programs in the following order. First, all variable assignments specified via the `-v` option are performed. Next, *gawk* compiles the program into an internal form. Then, *gawk* executes the code in the `BEGIN` rule(s) (if any), and then proceeds to read each file named in the `ARGV` array (up to `ARGV[ARGC-1]`). If there are no files named on the command line, *gawk* reads the standard input.

If a filename on the command line has the form `var=val` it is treated as a variable assignment. The variable `var` will be

assigned the value *val*. (This happens after any **BEGIN** rule(s) have been run.)

If the value of a particular element of **ARGV** is empty (""), *gawk* skips over it.

For each input file, if a **BEGINFILE** rule exists, *gawk* executes the associated code before processing the contents of the file. Similarly, *gawk* executes the code associated with **ENDFILE** rules after processing the file.

For each record in the input, *gawk* tests to see if it matches any *pattern* in the AWK program. For each pattern that the record matches, *gawk* executes the associated *action*. The patterns are tested in the order they occur in the program.

Finally, after all the input is exhausted, *gawk* executes the code in the **END** rule(s) (if any).

Command Line Directories

According to POSIX, files named on the *awk* command line must be text files. The behavior is “undefined” if they are not. Most versions of *awk* treat a directory on the command line as a fatal error.

For *gawk*, a directory on the command line produces a warning, but is otherwise skipped. If either of the **--posix** or **--traditional** options is given, then *gawk* reverts to treating directories on the command line as a fatal error.

VARIABLES, RECORDS AND FIELDS

[top](#)

AWK variables are dynamic; they come into existence when they are first used. Their values are either floating-point numbers or strings, or both, depending upon how they are used. Additionally, *gawk* allows variables to have regular-expression type. AWK also has one dimensional arrays; arrays with multiple dimensions may be simulated. However, *gawk* provides true arrays of arrays. Several pre-defined variables are set as a program runs; these are described as needed and summarized below.

Records

Normally, records are separated by newline characters. You can control how records are separated by assigning values to the built-in variable **RS**. See

https://www.gnu.org/software/gawk/manual/html_node/Records.html for the details.

Fields

As each input record is read, *gawk* splits the record into *fields*, using the value of the **FS** variable as the field separator.

Additionally, **FIELDWIDTHS** and **FPAT** may be used to control input field splitting. See the details, starting at

https://www.gnu.org/software/gawk/manual/html_node/Fields.html .

Each field in the input record may be referenced by its position: **\$1**, **\$2**, and so on. **\$0** is the whole record, including leading and trailing whitespace.

The variable **NF** is set to the total number of fields in the input record.

References to non-existent fields (i.e., fields after **\$NF**) produce the null string. However, assigning to a non-existent field (e.g., **\$(NF+2) = 5**) increases the value of **NF**, creates any intervening fields with the null string as their values, and causes the value of **\$0** to be recomputed, with the fields being separated by the value of **OFS**. References to negative numbered

fields cause a fatal error. Decrementing **NF** causes the values of fields past the new value to be lost, and the value of **\$0** to be recomputed, with the fields being separated by the value of **OFS**.

Assigning a value to an existing field causes the whole record to be rebuilt when **\$0** is referenced. Similarly, assigning a value to **\$0** causes the record to be resplit, creating new values for the fields.

Comma Separated Values

When invoked with either the **-k** or the **--csv** option, *gawk* does not use regular record determination and field splitting as described above. Instead, records are terminated by unquoted newlines, and fields are separated by commas. Double-quotes may be used to enclose fields containing commas, newlines, or doubled double-quotes. See https://www.gnu.org/software/gawk/manual/html_node/Comma-Separated-Fields.html for more details.

Built-in Variables

Gawk's built-in variables are listed below. This list is purposely terse. For details, see https://www.gnu.org/software/gawk/manual/html_node/Built_002din-Variables.

ARGC The number of command line arguments.

ARGIND The index in **ARGV** of the current file being processed.

ARGV Array of command line arguments. The array is indexed from 0 to **ARGC** - 1.

BINMODE

On non-POSIX systems, specifies use of "binary" mode for all file I/O. See https://www.gnu.org/software/gawk/manual/html_node/PC-Using.html for the details.

CONVFMT

The conversion format for numbers, "%.6g", by default.

ENVIRON

An array containing the values of the current environment. The array is indexed by the environment variables, each element being the value of that variable.

ERRNO If a system error occurs either doing a redirection for **getline**, during a read for **getline**, or during a **close()**, then **ERRNO** is set to a string describing the error. The value is subject to translation in non-English locales.

FIELDWIDTHS

A whitespace-separated list of field widths. When set, *gawk* parses the input into fields of fixed width, instead of using the value of the **FS** variable as the field separator. Each field width may optionally be preceded by a colon-separated value specifying the number of characters to skip before the field starts.

FILENAME

The name of the current input file. If no files are specified on the command line, the value of **FILENAME** is "-". However, **FILENAME** is undefined inside the **BEGIN** rule (unless set by **getline**).

FNR The input record number in the current input file.

FPAT A regular expression describing the contents of the fields in a record. When set, *gawk* parses the input into fields, where the fields match the regular expression, instead of

using the value of **FS** as the field separator.

FS The input field separator, a space by default. See https://www.gnu.org/software/gawk/manual/html_node/Field-Separators.html for the details.

FUNCTAB

An array whose indices and corresponding values are the names of all the user-defined or extension functions in the program. **NOTE:** You may not use the **delete** statement with the **FUNCTAB** array.

IGNORECASE

Controls the case-sensitivity of all regular expression and string operations. See https://www.gnu.org/software/gawk/manual/html_node/Case_002dsensitivity.html for details.

LINT Provides dynamic control of the **--lint** option from within an AWK program.

NF The number of fields in the current input record.

NR The total number of input records seen so far.

OFMT The output format for numbers, **"%.6g"**, by default.

OFS The output field separator, a space by default.

ORS The output record separator, by default a newline.

PREC The working precision of arbitrary precision floating-point numbers, 53 by default.

PROCINFO

The elements of this array provide access to information about the running AWK program. See https://www.gnu.org/software/gawk/manual/html_node/Auto_002dset.html for the details.

ROUNDMODE

The rounding mode to use for arbitrary precision arithmetic on numbers, by default **"N"** (IEEE-754 roundTiesToEven mode). See https://www.gnu.org/software/gawk/manual/html_node/Setting-the-rounding-mode.html for the details.

RS The input record separator, by default a newline.

RT The record terminator. *Gawk* sets **RT** to the input text that matched the character or regular expression specified by **RS**.

RSTART The index of the first character matched by **match()**; 0 if no match.

RLENGTH

The length of the string matched by **match()**; -1 if no match.

SUBSEP The string used to separate multiple subscripts in array elements, by default **"W034"**.

SYMTAB An array whose indices are the names of all currently defined global variables and arrays in the program. You may not use the **delete** statement with the **SYMTAB** array, nor assign to elements with an index that is not a variable name.

TEXTDOMAIN

The text domain of the AWK program; used to find the localized translations for the program's strings.

Arrays

Arrays are subscripted with an expression between square brackets ([and]). If the expression is an expression list (*expr*, *expr* ...) then the array subscript is a string consisting of the concatenation of the (string) value of each expression, separated by the value of the SUBSEP variable. This facility is used to simulate multiply dimensioned arrays. For example:

```
i = "A"; j = "B"; k = "C"
x[i, j, k] = "hello, worldWn"
```

assigns the string "hello, worldWn" to the element of the array x which is indexed by the string "AW034BW034C". All arrays in AWK are associative, i.e., indexed by string values.

The special operator `in` may be used to test if an array has an index consisting of a particular value:

```
if (val in array)
    print array[val]
```

If the array has multiple subscripts, use `(i, j) in array`.

The `in` construct may also be used in a `for` loop to iterate over all the elements of an array. However, the `(i, j) in array` construct only works in tests, not in `for` loops.

An element may be deleted from an array using the `delete` statement. The `delete` statement may also be used to delete the entire contents of an array, just by specifying the array name without a subscript.

gawk supports true multidimensional arrays. It does not require that such arrays be "rectangular" as in C or C++. See https://www.gnu.org/software/gawk/manual/html_node/Arrays for details.

Namespaces

Gawk provides a simple *namespace* facility to help work around the fact that all variables in AWK are global.

A *qualified name* consists of a two simple identifiers joined by a double colon (::). The left-hand identifier represents the namespace and the right-hand identifier is the variable within it. All simple (non-qualified) names are considered to be in the "current" namespace; the default namespace is `awk`. However, simple identifiers consisting solely of uppercase letters are forced into the `awk` namespace, even if the current namespace is different.

You change the current namespace with an `@namespace "name"` directive.

The standard predefined builtin function names may not be used as namespace names. The names of additional functions provided by *gawk* may be used as namespace names or as simple identifiers in other namespaces. For more details, see https://www.gnu.org/software/gawk/manual/html_node/Namespaces.html#Namespaces.

Variable Typing And Conversion

Variables and fields may be (floating point) numbers, or strings, or both. They may also be regular expressions. How the value of a variable is interpreted depends upon its context. If used in a numeric expression, it will be treated as a number; if used as a string it will be treated as a string.

To force a variable to be treated as a number, add zero to it; to force it to be treated as a string, concatenate it with the null string.

Uninitialized variables have the numeric value zero and the string value "" (the null, or empty, string).

When a string must be converted to a number, the conversion is accomplished using `strtod(3)`. A number is converted to a string by using the value of `CONVFMT` as a format string for `sprintf(3)`, with the numeric value of the variable as the argument. However, even though all numbers in AWK are floating-point, integral values are *always* converted as integers.

Gawk performs comparisons as follows: If two variables are numeric, they are compared numerically. If one value is numeric and the other has a string value that is a "numeric string," then comparisons are also done numerically. Otherwise, the numeric value is converted to a string and a string comparison is performed. Two strings are compared, of course, as strings.

Note that string constants, such as "57", are *not* numeric strings, they are string constants. The idea of "numeric string" only applies to fields, `getline` input, `FILENAME`, `ARGV` elements, `ENVIRON` elements and the elements of an array created by `split()` or `patsplit()` that are numeric strings. The basic idea is that *user input*, and only user input, that looks numeric, should be treated that way.

Octal and Hexadecimal Constants

You may use C-style octal and hexadecimal constants in your AWK program source code. For example, the octal value `011` is equal to decimal 9, and the hexadecimal value `0x11` is equal to decimal 17.

String Constants

String constants in AWK are sequences of characters enclosed between double quotes (like "value"). Within strings, certain *escape sequences* are recognized, as in C. See https://www.gnu.org/software/gawk/manual/html_node/Escape-Sequences for the details.

Regexp Constants

A regular expression constant is a sequence of characters enclosed between forward slashes (like `/value/`).

The escape sequences described in the manual may also be used inside constant regular expressions (e.g., `/[ WtWfWnWrWv]/` matches whitespace characters).

Gawk provides *strongly typed* regular expression constants. These are written with a leading `@` symbol (like so: `@/value/`). Such constants may be assigned to scalars (variables, array elements) and passed to user-defined functions. Variables that have been so assigned have regular expression type.

PATTERNS AND ACTIONS

[top](#)

AWK is a line-oriented language. The pattern comes first, and then the action. Action statements are enclosed in `{` and `}`. Either the pattern may be missing, or the action may be missing, but, of course, not both. If the pattern is missing, the action executes for every single record of input. A missing action is equivalent to

```
{ print }
```

which prints the entire record.

Comments begin with the `#` character, and continue until the end of the line. Empty lines may be used to separate statements. Normally, a statement ends with a newline, however, this is not the case for lines ending in a comma, `{`, `?`, `:`, `&&`, or `||`. Lines ending in `do` or `else` also have their statements automatically continued on the following line. In other cases, a line can be continued by ending it with a `"\`", in which case the newline is ignored. However, a `"\` after a `#` is not special.

Multiple statements may be put on one line by separating them with a `;`. This applies to both the statements within the action part of a pattern-action pair (the usual case), and to the pattern-action statements themselves.

Patterns

AWK patterns may be one of the following:

```
BEGIN
END
BEGINFILE
ENDFILE
/regular expression/
relational expression
pattern && pattern
pattern || pattern
pattern ? pattern : pattern
(pattern)
! pattern
pattern1, pattern2
```

BEGIN and **END** are two special kinds of patterns which are not tested against the input. The action parts of all **BEGIN** patterns are merged as if all the statements had been written in a single **BEGIN** rule. They are executed before any of the input is read. Similarly, all the **END** rules are merged, and executed when all the input is exhausted (or when an **exit** statement is executed). **BEGIN** and **END** patterns cannot be combined with other patterns in pattern expressions. **BEGIN** and **END** patterns cannot have missing action parts.

BEGINFILE and **ENDFILE** are additional special patterns whose actions are executed before reading the first record of each command-line input file and after reading the last record of each file. Inside the **BEGINFILE** rule, the value of **ERRNO** is the empty string if the file was opened successfully. Otherwise, there is some problem with the file and the code should use **nextfile** to skip it. If that is not done, *gawk* produces its usual fatal error for files that cannot be opened.

For */regular expression/* patterns, the associated statement is executed for each input record that matches the regular expression. Regular expressions are essentially the same as those in **egrep**(1). See https://www.gnu.org/software/gawk/manual/html_node/Regexp.html for the details on regular expressions.

A *relational expression* may use any of the operators defined below in the section on actions. These generally test whether certain fields match certain regular expressions.

The **&&**, **||**, and **!** operators are logical AND, logical OR, and logical NOT, respectively, as in C. They do short-circuit evaluation, also as in C, and are used for combining more primitive pattern expressions. As in most languages, parentheses may be used to change the order of evaluation.

The **?:** operator is like the same operator in C. If the first pattern is true then the pattern used for testing is the second

pattern, otherwise it is the third. Only one of the second and third patterns is evaluated.

The *pattern1*, *pattern2* form of an expression is called a *range pattern*. It matches all input records starting with a record that matches *pattern1*, and continuing until a record that matches *pattern2*, inclusive. It does not combine with any other sort of pattern expression.

Actions

Action statements are enclosed in braces, { and }. Action statements consist of the usual assignment, conditional, and looping statements found in most languages. The operators, control statements, and input/output statements available are patterned after those in C.

Operators

The operators in AWK, in order of decreasing precedence, are:

(...) Grouping

\$ Field reference.

++ -- Increment and decrement, both prefix and postfix.

^ Exponentiation.

+ - ! Unary plus, unary minus, and logical negation.

* / % Multiplication, division, and modulus.

+ - Addition and subtraction.

space String concatenation.

| |& Piped I/O for `getline`, `print`, and `printf`.

< > <= >= == !=
The regular relational operators.

~ !~ Regular expression match, negated match.

in Array membership.

&& Logical AND.

|| Logical OR.

?: The C conditional expression. This has the form *expr1* ? *expr2* : *expr3*. If *expr1* is true, the value of the expression is *expr2*, otherwise it is *expr3*. Only one of *expr2* and *expr3* is evaluated.

= += -= *= /= %= ^=
Assignment. Both absolute assignment (*var* = *value*) and operator-assignment (the other forms) are supported.

Control Statements

The control statements are as follows:

```
if (condition) statement [ else statement ]
while (condition) statement
do statement while (condition)
for (expr1; expr2; expr3) statement
for (var in array) statement
break
continue
delete array[index]
delete array
```

```

exit [ expression ]
{ statements }
switch (expression) {
case value|regex : statement
...
[ default: statement ]
}

```

I/O Statements

The input/output statements are as follows:

close(*file* [, *how*])

Close an open file, pipe or coprocess. The optional *how* should only be used when closing one end of a two-way pipe to a coprocess. It must be a string value, either "to" or "from".

getline

Set \$0 from the next input record; set NF, NR, FNR, RT.

getline <*file*

Set \$0 from the next record of *file*; set NF, RT.

getline *var*

Set *var* from the next input record; set NR, FNR, RT.

getline *var* <*file*

Set *var* from the next record of *file*; set RT.

command | **getline** [*var*]

Run *command*, piping the output either into \$0 or *var*, as above, and RT.

command |& **getline** [*var*]

Run *command* as a coprocess piping the output either into \$0 or *var*, as above, and RT. (The *command* can also be a socket. See the subsection **Special File Names**, below.)

fflush([*file*])

Flush any buffers associated with the open output file or pipe *file*. If *file* is missing or if it is the null string, then flush all open output files and pipes.

next Stop processing the current input record. Read the next input record and start processing over with the first pattern in the AWK program. Upon reaching the end of the input data, execute any **END** rule(s).

nextfile

Stop processing the current input file. The next input record read comes from the next input file. Update **FILENAME** and **ARGIND**, reset **FNR** to 1, and start processing over with the first pattern in the AWK program. Upon reaching the end of the input data, execute any **ENDFILE** and **END** rule(s).

print Print the current record. The output record is terminated with the value of **ORS**.

print *expr-list*

Print expressions. Each expression is separated by the value of **OFS**. The output record is terminated with the value of **ORS**.

print *expr-list* >*file*

Print expressions on *file*. Each expression is separated by the value of **OFS**. The output record is terminated with the value of **ORS**.

```
printf fmt, expr-list
    Format and print.
```

```
printf fmt, expr-list >file
    Format and print on file.
```

```
system(cmd-line)
    Execute the command cmd-line, and return the exit status.
    (This may not be available on non-POSIX systems.) See
    https://www.gnu.org/software/gawk/manual/html\_node/\_I\_002f0-Functions.html#\_I\_002f0-Functions
    for the full details on the exit status.
```

Additional output redirections are allowed for `print` and `printf`.

```
print ... >> file
    Append output to the file.
```

```
print ... | command
    Write on a pipe.
```

```
print ... |& command
    Send data to a coprocess or socket. (See also the
    subsection Special File Names, below.)
```

The `getline` function returns 1 on success, zero on end of file, and -1 on an error. If the `errno(3)` value indicates that the I/O operation may be retried, and `PROCINFO["input", "RETRY"]` is set, then -2 is returned instead of -1, and further calls to `getline` may be attempted. Upon an error, `ERRNO` is set to a string describing the problem.

NOTE: Failure in opening a two-way socket results in a non-fatal error being returned to the calling function. If using a pipe, coprocess, or socket to `getline`, or from `print` or `printf` within a loop, you *must* use `close()` to create new instances of the command or socket. AWK does not automatically close pipes, sockets, or coprocesses when they return EOF.

The AWK versions of the `printf` statement and `sprintf()` function are similar to those of C. For details, see https://www.gnu.org/software/gawk/manual/html_node/Printf.html.

Special File Names

When doing I/O redirection from either `print` or `printf` into a file, or via `getline` from a file, *gawk* recognizes certain special filenames internally. These filenames allow access to open file descriptors inherited from *gawk*'s parent process (usually the shell). These file names may also be used on the command line to name data files. The filenames are:

- The standard input.

```
/dev/stdin
    The standard input.
```

```
/dev/stdout
    The standard output.
```

```
/dev/stderr
    The standard error output.
```

```
/dev/fd/n
    The file associated with the open file descriptor n.
```

The following special filenames may be used with the `|&` coprocess operator for creating TCP/IP network connections:

```
/inet/tcp/lport/rhost/rport
/inet4/tcp/lport/rhost/rport
```

`/inet6/tcp/ lport / rhost / rport`

Files for a TCP/IP connection on local port *lport* to remote host *rhost* on remote port *rport*. Use a port of 0 to have the system pick a port. Use `/inet4` to force an IPv4 connection, and `/inet6` to force an IPv6 connection. Plain `/inet` uses the system default (most likely IPv4). Usable only with the `|&` two-way I/O operator.

`/inet/udp/ lport / rhost / rport`

`/inet4/udp/ lport / rhost / rport`

`/inet6/udp/ lport / rhost / rport`

Similar, but use UDP/IP instead of TCP/IP. UDP support is obsolete and will be removed in the future.

Numeric Functions

AWK has the following built-in arithmetic functions:

`atan2(y, x)`

Return the arctangent of *y/x* in radians.

`cos(expr)`

Return the cosine of *expr*, which is in radians.

`exp(expr)`

The exponential function.

`int(expr)`

Truncate to integer.

`log(expr)`

The natural logarithm function.

`rand()` Return a random number *N*, between zero and one, such that $0 \leq N < 1$.

`sin(expr)`

Return the sine of *expr*, which is in radians.

`sqrt(expr)`

Return the square root of *expr*.

`srand([expr])`

Use *expr* as the new seed for the random number generator. If no *expr* is provided, use the time of day. Return the previous seed for the random number generator.

String Functions

Gawk has the following built-in string functions; details are provided in

https://www.gnu.org/software/gawk/manual/html_node/String-Functions .

`asort(s [, d [, how ]])`

Return the number of elements in the source array *s*. Sort the contents of *s* using *gawk*'s normal rules for comparing values, and replace the indices of the sorted values *s* with sequential integers starting with 1. If the optional destination array *d* is specified, first duplicate *s* into *d*, and then sort *d*, leaving the indices of the source array *s* unchanged. The optional string *how* controls the direction and the comparison mode. Valid values for *how* are described in

https://www.gnu.org/software/gawk/manual/html_node/String-Functions.html#String-Functions .

s and *d* are allowed to be the same array; this only makes sense when supplying the third argument as well.

`asorti(s [, d [, how ]])`

Return the number of elements in the source array *s*. The behavior is the same as that of `asort()`, except that the array *indices* are used for sorting, not the array values.

When done, the array is indexed numerically, and the values are those of the original indices. The original values are lost; thus provide a second array if you wish to preserve the original. The purpose of the optional string *how* is the same as for `asort()`. Here too, *s* and *d* are allowed to be the same array; this only makes sense when supplying the third argument as well.

`gensub(r, s, h [, t])`

Search the target string *t* for matches of the regular expression *r*. If *h* is a string beginning with **g** or **G**, then replace all matches of *r* with *s*. Otherwise, *h* is a number indicating which match of *r* to replace. If *t* is not supplied, use `$0` instead. Within the replacement text *s*, the sequence `Wn`, where *n* is a digit from 1 to 9, may be used to indicate just the text that matched the *n*'th parenthesized subexpression. The sequence `W0` represents the entire matched text, as does the character `&`. Unlike `sub()` and `gsub()`, the modified string is returned as the result of the function, and the original target string is *not* changed.

`gsub(r, s [, t])`

For each substring matching the regular expression *r* in the string *t*, substitute the string *s*, and return the number of substitutions. If *t* is not supplied, use `$0`. An `&` in the replacement text is replaced with the text that was actually matched. Use `W&` to get a literal `&`. (This must be typed as `"WW&"`; see

https://www.gnu.org/software/gawk/manual/html_node/Gory-Details.html#Gory-Details for a fuller discussion of the rules for ampersands and backslashes in the replacement text of `sub()`, `gsub()`, and `gensub()`.)

`index(s, t)`

Return the index of the string *t* in the string *s*, or zero if *t* is not present. (This implies that character indices start at one.)

`length([s])`

Return the length of the string *s*, or the length of `$0` if *s* is not supplied. With an array argument, `length()` returns the number of elements in the array.

`match(s, r [, a])`

Return the position in *s* where the regular expression *r* occurs, or zero if *r* is not present, and set the values of `RSTART` and `RLENGTH`. Note that the argument order is the same as for the `~` operator: *str* `~` *re*. See

https://www.gnu.org/software/gawk/manual/html_node/String-Functions.html#String-Functions for a description of how the array *a* is filled if it is provided.

`patsplit(s, a [, r [, seps] ])`

Split the string *s* into the array *a* and the separators array *seps* on the regular expression *r*, and return the number of fields. Element values are the portions of *s* that matched *r*. The value of *seps*[*i*] is the possibly null separator that appeared after *a*[*i*]. The value of *seps*[0] is the possibly null leading separator. If *r* is omitted, `FPAT` is used instead. The arrays *a* and *seps* are cleared first. Splitting behaves identically to field splitting with `FPAT`.

`split(s, a [, r [, seps] ])`

Split the string *s* into the array *a* and the separators array *seps* on the regular expression *r*, and return the number of fields. If *r* is omitted, `FS` is used instead. The arrays *a* and *seps* are cleared first. *seps*[*i*] is the

field separator matched by *r* between *a[i]* and *a[i+1]*.
Splitting behaves identically to field splitting.

sprintf(*fmt*, *expr-list*)

Print *expr-list* according to *fmt*, and return the resulting string.

strtonum(*str*)

Examine *str*, and return its numeric value. If *str* begins with a leading 0, treat it as an octal number. If *str* begins with a leading 0x or 0X, treat it as a hexadecimal number. Otherwise, assume it is a decimal number.

sub(*r*, *s* [, *t*])

Just like **gsub()**, but replace only the first matching substring. Return either zero or one.

substr(*s*, *i* [, *n*])

Return the at most *n*-character substring of *s* starting at *i*. If *n* is omitted, use the rest of *s*.

tolower(*str*)

Return a copy of the string *str*, with all the uppercase characters in *str* translated to their corresponding lowercase counterparts. Non-alphabetic characters are left unchanged.

toupper(*str*)

Return a copy of the string *str*, with all the lowercase characters in *str* translated to their corresponding uppercase counterparts. Non-alphabetic characters are left unchanged.

Gawk is multibyte aware. This means that **index()**, **length()**, **substr()** and **match()** all work in terms of characters, not bytes.

Time Functions

Gawk provides the following functions for obtaining time stamps and formatting them. Details are provided in https://www.gnu.org/software/gawk/manual/html_node/Time-Functions.

mktime(*datespec* [, *utc-flag*])

Turn *datespec* into a time stamp of the same form as returned by **systeme()**, and return the result. If *utc-flag* is present and is non-zero or non-null, the time is assumed to be in the UTC time zone; otherwise, the time is assumed to be in the local time zone. If *datespec* does not contain enough elements or if the resulting time is out of range, **mktime()** returns -1. See https://www.gnu.org/software/gawk/manual/html_node/Time-Functions.html#Time-Functions for the details of *datespec*.

strftime([*format* [, *timestamp* [, *utc-flag*]])

Format *timestamp* according to the specification in *format*. If *utc-flag* is present and is non-zero or non-null, the result is in UTC, otherwise the result is in local time. The *timestamp* should be of the same form as returned by **systeme()**. If *timestamp* is missing, the current time of day is used. If *format* is missing, a default format equivalent to the output of **date(1)** is used. The default format is available in **PROCINFO["strftime"]**. See the specification for the **strftime()** function in ISO C for the format conversions that are guaranteed to be available.

systeme()

Return the current time of day as the number of seconds since the Epoch (1970-01-01 00:00:00 UTC on POSIX systems).

Bit Manipulations Functions

Gawk supplies the following bit manipulation functions. They work by converting double-precision floating point values to `uintmax_t` integers, doing the operation, and then converting the result back to floating point. Passing negative operands to any of these functions causes a fatal error.

The functions are:

`and(v1, v2 [, ...])`

Return the bitwise AND of the values provided in the argument list. There must be at least two.

`compl(val)`

Return the bitwise complement of *val*.

`lshift(val, count)`

Return the value of *val*, shifted left by *count* bits.

`or(v1, v2 [, ...])`

Return the bitwise OR of the values provided in the argument list. There must be at least two.

`rshift(val, count)`

Return the value of *val*, shifted right by *count* bits.

`xor(v1, v2 [, ...])`

Return the bitwise XOR of the values provided in the argument list. There must be at least two.

Type Functions

The following functions provide type related information about their arguments.

`isarray(x)`

Return true if *x* is an array, false otherwise.

`typeof(x)`

Return a string indicating the type of *x*. The string will be one of "array", "number", "number|bool", "regexp", "string", "strnum", "unassigned", or "untyped".

Internationalization Functions

The following functions may be used from within your AWK program for translating strings at run-time. For full details, see

https://www.gnu.org/software/gawk/manual/html_node/18N-Functions.html#18N-Functions .

`bindtextdomain(directory [, domain])`

Specify the directory where *gawk* looks for the `.gmo` files, in case they will not or cannot be placed in the "standard" locations. It returns the directory where *domain* is "bound."

The default *domain* is the value of `TEXTDOMAIN`. If *directory* is the null string (""), then `bindtextdomain()` returns the current binding for the given *domain*.

`dcgettext(string [, domain [, category]])`

Return the translation of *string* in text domain *domain* for locale category *category*. The default value for *domain* is the current value of `TEXTDOMAIN`. The default value for *category* is "LC_MESSAGES".

`dcngettext(string1, string2, number [, domain [, category]])`

Return the plural form used for *number* of the translation of *string1* and *string2* in text domain *domain* for locale category *category*. The default value for *domain* is the current value of `TEXTDOMAIN`. The default value for *category* is "LC_MESSAGES".

Boolean Valued Functions

You can create special Boolean-typed values; see the manual for how they work and why they exist.

`mkbool(expression)`

Based on the boolean value of *expression* return either a true value or a false value. True values have numeric value one. False values have numeric value zero.

USER-DEFINED FUNCTIONS

[top](#)

Functions in AWK are defined as follows:

```
function name(parameter list) { statements }
```

Functions execute when they are called from within expressions in either patterns or actions. Actual parameters supplied in the function call are used to instantiate the formal parameters declared in the function. Arrays are passed by reference, other variables are passed by value.

Local variables are declared as extra parameters in the parameter list. The convention is to separate local variables from real parameters by extra spaces in the parameter list. For example:

```
function f(p, q,      a, b)  # a and b are local
{
    ...
}

/abc/ { ... ; f(1, 2) ; ... }
```

The left parenthesis in a function call is required to immediately follow the function name, without any intervening whitespace. This restriction does not apply to the built-in functions listed above.

Functions may call each other and may be recursive. Function parameters used as local variables are initialized to the null string and the number zero upon function invocation.

Use `return expr` to return a value from a function. The return value is undefined if no value is provided, or if the function returns by “falling off” the end.

Functions may be called indirectly. To do this, assign the name of the function to be called, as a string, to a variable. Then use the variable as if it were the name of a function, prefixed with an `@` sign, like so:

```
function myfunc()
{
    print "myfunc called"
    ...
}

{
    ...
    the_func = "myfunc"
    @the_func()  # call through the_func to myfunc
    ...
}
```

If `--lint` has been provided, *gawk* warns about calls to undefined functions at parse time, instead of at run time. Calling an undefined function at run time is a fatal error.

DYNAMICALLY LOADING NEW FUNCTIONS

[top](#)

You can dynamically add new functions written in C or C++ to the running *gawk* interpreter with the `@load` statement. The full details are beyond the scope of this manual page; see https://www.gnu.org/software/gawk/manual/html_node/Dynamic-Extensions.html#Dynamic-Extensions .

SIGNALS [top](#)

The *gawk* profiler accepts two signals. `SIGUSR1` causes it to dump a profile and function call stack to the profile file, which is either `awkprof.out`, or whatever file was named with the `--profile` option. It then continues to run. `SIGHUP` causes *gawk* to dump the profile and function call stack and then exit.

INTERNATIONALIZATION [top](#)

String constants are sequences of characters enclosed in double quotes. In non-English speaking environments, it is possible to mark strings in the AWK program as requiring translation to the local natural language. Such strings are marked in the AWK program with a leading underscore (`"_"`). For example,

```
gawk 'BEGIN { print "hello, world" }'
```

always prints `hello, world`. But,

```
gawk 'BEGIN { print _"hello, world" }'
```

might print `bonjour, monde` in France. See https://www.gnu.org/software/gawk/manual/html_node/Internationalization.html#Internationalization for the steps involved in producing and running a localizable AWK program.

GNU EXTENSIONS [top](#)

Gawk has a too-large number of extensions to POSIX *awk*. They are described in https://www.gnu.org/software/gawk/manual/html_node/POSIX_002fGNU.html . All the extensions can be disabled by invoking *gawk* with the `--traditional` or `--posix` options.

ENVIRONMENT VARIABLES [top](#)

The `AWKPATH` environment variable can be used to provide a list of directories that *gawk* searches when looking for files named via the `-f`, `--file`, `-i` and `--include` options, and the `@include` and `@nsinclude` directives. If the initial search fails, the path is searched again after appending `.awk` to the filename.

The `AWKLIBPATH` environment variable can be used to provide a list of directories that *gawk* searches when looking for files named via the `-l` and `--load` options.

The `GAWK_PERSIST_FILE` environment variable, if present, specifies a file to use as the backing store for persistent memory. See *GAWK: Effective AWK Programming* for the details.

The `GAWK_READ_TIMEOUT` environment variable can be used to specify a timeout in milliseconds for reading input from a terminal, pipe or two-way communication including sockets.

For connection to a remote host via socket, `GAWK SOCK RETRIES` controls the number of retries, and `GAWK MSEC SLEEP` the interval between retries. The interval is in milliseconds. On systems that do not support `usleep(3)`, the value is rounded up to an

integral number of seconds.

If `POSIXLY_CORRECT` exists in the environment, then *gawk* behaves exactly as if `--posix` had been specified on the command line. If `--lint` has been specified, *gawk* issues a warning message to this effect.

EXIT STATUS [top](#)

If the `exit` statement is used with a value, then *gawk* exits with the numeric value given to it.

Otherwise, if there were no problems during execution, *gawk* exits with the value of the C constant `EXIT_SUCCESS`. This is usually zero.

If an error occurs, *gawk* exits with the value of the C constant `EXIT_FAILURE`. This is usually one.

If *gawk* exits because of a fatal error, the exit status is 2. On non-POSIX systems, this value may be mapped to `EXIT_FAILURE`.

VERSION INFORMATION [top](#)

This man page documents *gawk*, version 5.4.

AUTHORS [top](#)

The original version of UNIX *awk* was designed and implemented by Alfred Aho, Peter Weinberger, and Brian Kernighan of Bell Laboratories. Ozan Yigit is the current maintainer. Brian Kernighan occasionally dabbles in its development.

Paul Rubin and Jay Fenlason, of the Free Software Foundation, wrote *gawk*, to be compatible with the original version of *awk* distributed in Seventh Edition UNIX. John Woods contributed a number of bug fixes. David Trueman, with contributions from Arnold Robbins, made *gawk* compatible with the new version of UNIX *awk*. Arnold Robbins is the current maintainer.

See *GAWK: Effective AWK Programming* for a full list of the contributors to *gawk* and its documentation.

See the `README` file in the *gawk* distribution for up-to-date information about maintainers and which ports are currently supported.

BUG REPORTS AND QUESTIONS [top](#)

If you find a bug in *gawk*, please use the `gawkbug(1)` program to report it.

Full instructions for reporting a bug are provided in https://www.gnu.org/software/gawk/manual/html_node/Bugs.html. Please carefully read and follow the instructions given there. This will make bug reporting and resolution much easier for everyone involved. Really.

On the other hand, if you have a question as to how to accomplish a particular task using *awk* or *gawk*, send an email to `help-gawk@gnu.org` with your request for help.

BUGS [top](#)

The `-F` option is not necessary given the command line variable assignment feature; it remains only for backwards compatibility.

This manual page is too long: *gawk* has too many features.

SEE ALSO [top](#)

[egrep\(1\)](#), [sed\(1\)](#), [gawkbug\(1\)](#), [printf\(3\)](#), and [strftime\(3\)](#).

The AWK Programming Language, second edition, Alfred V. Aho, Brian W. Kernighan, Peter J. Weinberger, Addison-Wesley, 2023. ISBN 9-780138-269722.

GAWK: Effective AWK Programming, Edition 5.4, shipped with the *gawk* source. The current version of this document is available online at <https://www.gnu.org/software/gawk/manual>.

The GNU `gettext` documentation, available online at <https://www.gnu.org/software/gettext>.

EXAMPLES [top](#)

Print and sort the login names of all users:

```
BEGIN    { FS = ":" }
          { print $1 | "sort" }
```

Count lines in a file:

```
          { nlines++ }
END      { print nlines }
```

Precede each line by its number in the file:

```
{ print FNR, $0 }
```

Concatenate and line number (a variation on a theme):

```
{ print NR, $0 }
```

Run an external command for particular lines of data:

```
tail -f access_log |
awk '/myhome.html/ { system("nmap " $1 ">> logdir/myhome.html") }'
```

COPYING PERMISSIONS [top](#)

Copyright © 1989, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2001, 2002, 2003, 2004, 2005, 2007, 2009, 2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025, 2026 Free Software Foundation, Inc.

Permission is granted to make and distribute verbatim copies of this manual page provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual page under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual page into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the Foundation.

COLOPHON[top](#)

This page is part of the *gawk* (GNU awk) project. Information about the project can be found at <http://www.gnu.org/software/gawk/>. If you have a bug report for this manual page, see <http://pkg-shadow.alioth.debian.org/getinvolved.php>. This page was obtained from the project's upstream Git repository <git://git.savannah.gnu.org/gawk.git> on 2026-05-24. (At that time, the date of the most recent commit that was found in the repository was 2026-05-17.) If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date source for the page, or you have corrections or improvements to the information in this COLOPHON (which is *not* part of the original manual page), send a mail to man-pages@man7.org

Free Software Foundation

March 2 2026

GAWK(1)

HTML rendering created 2026-05-30 by [Michael Kerrisk](#), author of *The Linux Programming Interface*.

For details of in-depth **Linux/UNIX system programming training courses** that I teach, look [here](#).

Hosting by [jambit GmbH](#).

